

# Reconfiguration of Spanning Trees in Networks in the Presence of Node Failures

K. Ravindran, G. Singh and P. Gupta

Department of Computing and Information Sciences  
Kansas State University  
Manhattan, KS 66506

**Abstract:** Connectivity among a set of user entities in a network can be provided by a network level abstraction of acyclic graph (or spanning tree). The paper discusses the reconfiguration of a graph in the presence of failure of network nodes. A reconfiguration manifests as a *graph fragmentation* problem, whereby two or more disjoint subgraphs attempt to connect with one another to form a composite graph. Fragment interconnection requires contention resolution between fragments to avoid cycles. This paper presents two classes of contention resolution algorithms applicable for environments with potentially large number of fragments. They are based on pre-established ranking among fragments and random arbitration among fragments. The algorithms have been evaluated by simulation and compared. The algorithms are useful in supporting data multicasting across workstations and distributed computations involving data on different machines.

## 1 Introduction

One of the research problems on wide-area networks focuses on providing connectivity among users entities (or simply users) attached to various network nodes. The connectivity can be best provided by a network level abstraction of *acyclic graph* (or spanning tree) that has vertices in network nodes and edges over communication links between the nodes, with users residing in one or more vertices. All vertices and edges in a graph can participate in the flow of information across users. Typical applications are data multicasting across workstations in wide-area networks [1, 2], finding routes between transport hosts in a wide-area network [3, 4] and large scale parallel computations involving functions placed on different processors in a interconnected processor engine (*e.g.*, hypercubes).

A graph represents a logical topology of vertices and edges superimposed on the physical topology of nodes and links. This paper concerns with reconfiguration of a network graph, i.e., at the logical topology level, caused by node removals from the network due to failures (say, a crash). In this reconfiguration, the failure of a node containing a vertex causes execution of an algorithm by other nodes in the network that creates edges in the graph

through rest of the network to circumvent the failed node and maintain connectivity among the remaining vertices. We assume that the network remains connected at the physical connectivity level in the presence of failures.

Typically, the failure of a node containing a vertex partitions the graph into one or more subgraphs not connected with one another, which we refer to as *fragments*. We model the reconfiguration problem as one of connecting two or more fragments together by creating edges. A reconfiguration algorithm is a distributed protocol executed by each fragment to connect to the other fragments. The fragment interconnection requires coordination among the concerned fragments so that the resulting graph composed from the fragments does not contain cycles. The coordination problem manifests itself as resolving contention among fragments attempting creation of edges simultaneously to determine which edges should be created and which should not be. The paper discusses the reconfiguration problem from this perspective of fragment coordination. The paper describes two classes of contention resolution protocols that are usable in an environment where the number of fragments can be potentially large (for instance, high speed network applications involving a large number of users):

- Based on a pre-established ranking among the fragments;
- Based on a random arbitration among the fragments.

The ranking based protocol results in a deterministic algorithm (i.e., an algorithm that terminates with known delay bounds). However, the algorithm needs to derive the ranking of fragments from node identifiers (id) in the underlying network. The random arbitration protocol results in a probabilistic algorithm. Here, the algorithm needs only a function to generate random numbers that allows the contending fragments to re-attempt interconnection after a random interval of time, but otherwise does not need any special support mechanism from the network. The performance of the two types of algorithms have been studied by simulation.

The algorithm deals with the case when there are no failures during a reconfiguration. However, the algorithm can be used as a building block in composing an extended




Figure 1: Graph abstractions in network

algorithm to handle the case of such failures. The algorithmic framework of interconnecting graph fragments can also be generalized into a reconfiguration algorithm that deals with vertices being added and/or removed from a graph, whereby a node joining or leaving the network can also participate in the algorithm [6].

The paper is organized as follows: Section 2 describes the model of a graph and the concept of fragmentation and reconnection. Section 3 describes a base algorithm. Section 4 describes the ranking based contention resolution. Section 5 describes the random arbitration based contention resolution. Section 6 evaluates the two classes of algorithms. Section 7 summarizes the paper.

## 2 Reconfiguration Model

The network consists of a set of nodes  $\mathcal{V}$  interconnected through point-to-point links (e.g., data switches linked by fiber optic cables in a high speed network, processors linked by interconnection bus in a hypercube). The user entities implementing an application may be attached to one or more network nodes  $\mathcal{U}$ , and are end-points of communication. The network has sufficient connectivity so that every node is reachable from every other node. See Figure 1.

A graph  $\mathcal{G}(\hat{\mathcal{U}}, \hat{\mathcal{E}})$  provides a communication path (i.e., logical connectivity) among the users through a set of network nodes  $\hat{\mathcal{U}}$ , where  $\mathcal{U} \subseteq \hat{\mathcal{U}} \subseteq \mathcal{V}$  and edges  $\hat{\mathcal{E}}$  map one-to-one to the intervening links.  $\mathcal{G}$  may be viewed as an unrooted spanning tree for users, whereby the user residing in a node  $u \in \mathcal{U}$  may execute algorithms for traversing a tree  $\mathcal{T}_u(\mathcal{G})$  with root at  $u$  and leaves at the nodes  $(\mathcal{U} - u)$ . For instance, the user may be a transport level host in a network and the tree traversal over the edges may allow identification of data routes to the other hosts [4].  $\mathcal{G}$  thus embeds a set of trees  $\{\mathcal{T}_u(\mathcal{G})\}_{\forall u \in \mathcal{U}}$ .

When a node containing a vertex  $u'$  fails (say, due to a crash), a reconfiguration algorithm in the network creates a set of edges to interconnect the various vertices  $\in \hat{\mathcal{U}}$  that are adjacent to  $u'$ . Such an algorithm basically causes re-computing the data paths in  $\mathcal{G}$  through the newly created edges.

### 2.1 Graph fragmentation

A fundamental problem in logical topology reconfigurations is the handling of *fragmentation* in a graph, whereby two or more subgraphs not connected with one another (or fragments) need to be connected together by creating edges and vertices over appropriate links and nodes respectively in the physical topology to form a composite graph. Typically, the failure of a node containing a vertex fragments the graph at this vertex so that these fragments reconnect with one another circumventing the failed node. See figure 2. Prior to failure, node  $S_9$  had graph edges to its neighbors  $S_6, S_7, S_8, S_{10}$  and  $S_{11}$ . After the failure, the sets of nodes  $\{S_1, S_3, S_4, S_7\}$ ,  $\{S_2, S_5, S_6\}$ ,  $\{S_{10}, S_{12}\}$ ,  $\{S_8\}$  and  $\{S_{11}\}$  constitute the fragments.

### 2.2 Fragment Interconnection

A fragment is identified by the id of one of its constituent nodes, namely a neighbor of the failed node that has a vertex and detects the failure (e.g., in Figure 2,  $\{S_{10}, S_{12}\}$  is identified by the node id of  $S_{10}$ ). The node which detects the failure starts the reconfiguration activity on behalf of its fragment to find physical paths to connect to other fragments. Often, there can be several neighbors of a failed node which may detect the failure and each of these nodes may initiate a reconfiguration activity.

When a set of fragments attempt to connect with one another, it is necessary to prevent cycles from being formed in the composite graph as a result of the reconfiguration. Preventing of cycles basically requires the fragments to coordinate with one another in some form. Consider, for instance, two fragments  $F_1$  and  $F_2$  attempting reconfiguration. If  $F_1$  and  $F_2$  connect independently with each other, the resulting graph can contain cycles. To prevent this, a possibility is for  $F_1$  and  $F_2$  to coordinate with each other so that only one of them connects to the other. We model this coordination problem as one of resolving contention among the fragments to elect a coordinator which then creates edges to interconnect the various fragments. Refer to Figure 2. Suppose the fragment  $\{S_{10}, S_{12}\}$  becomes the coordinator after contention with other fragments. The nodes  $S_{10}$  and  $S_{12}$  execute a protocol to create edges between the pairs of nodes (say)  $[S_4, S_6]$ ,  $[S_7, S_8]$ ,  $[S_6, S_{10}]$  and  $[S_{11}, S_{12}]$ .

Thus, there are two basic components of a reconfiguration algorithm, viz., the election and the creation of edges. A simple but less efficient method may be to execute these components as distinct activities in that sequence. In our approach however, we have chosen to integrate these components into the algorithm as inseparable activities. This is to reduce the message overhead and the reconfiguration latency, i.e., time for reconfiguration to complete. In this approach, each fragment attempts to grow in size by creating edges towards other fragments until it meets another fragment. A contention between the fragments results at the point of meeting, whereupon one of the fragments wins the contest (using some contention resolution mechanism).

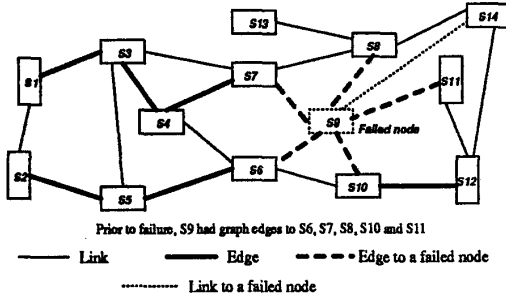


Figure 2: Fragmentation of graph

The winning fragment then expands by acquiring the losing fragment and attempts to further grow in size. The algorithm eventually completes by generating a single winner, with all fragments in the reconfiguration set connected together as a composite fragment acquired by the winner.

### 3 A base level algorithm

This section discusses the various aspects of our algorithmic realization of a reconfiguration activity.

#### 3.1 Network model

The physical connectivity information is maintained in a node internal table. Consider a node  $S$ . Let  $\{S^j\}_{j=1,2,\dots,K}$  be a set of nodes that are neighbors to  $S$ , i.e., nodes with a link from  $S$ . Then  $\{(id(S^j), port_j)\}$  constitutes the physical connectivity table of  $S$ , where  $id(S^j)$  is the networkwide unique id of  $S^j$  and  $port_j$  refers to a node level port (i.e., network interface) through which  $S$  is linked to  $S^j$ . The network implements a function **Port.To**( $S^j$ ) to extract  $port_j$  from the table and a function **Set.of.ports** to extract the list of ports to all neighboring nodes. With a unique id for every node in the network, the tables at all nodes specify the physical connectivity in the network.  $S$  may execute a separate protocol whereby  $S$  can detect the removal of a neighbor (say due to failure) and the addition of a node as a neighbor, and update the table accordingly.

A function **Assign\_Edge**( $port_j$ ) is provided that creates a directed edge from  $S$  originating through  $port_j$  and incident on the neighbor  $S^j$  so that  $S$  becomes adjacent to  $S^j$  in the graph. A function **Get\_Port** is also provided that generates a list of ports through which  $S$  has edges to its adjoining vertices.

We assume that the network provides reliable FIFO transmission of messages over a link and that transmission delay over a link is bounded. We also assume that the network diameter  $D$  is finite.

#### 3.2 State variables

The information structures maintained by the algorithm are distributed, so each vertex only knows its adjacent vertices. In this sense, we provide a decentralised structure of reconfiguration algorithm. In our algorithm, each node maintains the following variables:

1. *Coord.so\_far*: This is the local view of a node as to which fragment is coordinating. When a node detects the failure of a neighbor and starts a reconfiguration, it sets *Coord.so\_far* to its own id *my\_id*.
2. *Port.to\_coord*: This is the port in a node through which the fragment identified in *Coord.so\_far* can be reached. When *Coord.so\_far* = *my\_id*, the *Port.to\_coord* is set to *NULL*.
3. *status*: This indicates either IDLE or WAIT, meaning respectively that a node is not participating or otherwise in a reconfiguration.
4. *recd\_reply*[ $\cdot$ ]: This gives information about a port as to whether a response has been received through that port indicating completion of the reconfiguration activity in that direction or otherwise.

#### 3.3 Propagation of reconfiguration information

A node  $S$  sends a **Reconfig**(*node\_list*, *frag\_id*) message through one or more of its ports, where *node\_list* is a list of nodes — sorted in last-in-first-out order — through which the message has already traversed, and *frag\_id* refers to the id of a neighbor that detected the failure of node. If  $S$  is the initiator of the reconfiguration, *node\_list* and *frag\_id* are each set to *my\_id*, and **Reconfig** is sent through all the ports. Otherwise, the **Reconfig** is triggered by the arrival of a **Reconfig** at  $S$  from a neighbor; in this case, *my\_id* is appended to *node\_list*, and the **Reconfig** is forwarded through all ports except the port through which the **Reconfig** arrived. The flooding of messages is to propagate the information through the network that *frag\_id* contends to become the coordinator for reconfiguration.

A node  $S'$  is initially in IDLE state. Upon receiving a **Reconfig** message,  $S'$  node sets *Coord.so\_far* = *frag\_id* and *Port.to\_coord* = **Port.To**(*sender\_of*(**Reconfig**)) and moves into WAIT state, awaiting completion of the reconfiguration in the downstream direction (i.e., away from the coordinator).

$S'$  generates a response message to *sender\_of*(**Reconfig**) when there is no contention in the downstream direction from  $S'$ . The response message, referred to as **No\_contention**, implies that a new edge cannot be created along the path it traverses. The message is generated by default in the case when  $S'$  does not have a port to forward the **Reconfig**, i.e., **Set.of.ports** = {**Port.To**(*sender\_of*(**Reconfig**))} +  $Q$ , where  $Q$  is the set of ports (if any) along which either fragments have already been acquired or contention does not exist. The above actions are given by the following code skeleton:

```

receive Reconfig(node_list, frag_id)
if (status = IDLE)
    Generate No_contention if no more ports
    to forward Reconfig
    Coord_so_far := frag_id;
    status := WAIT;
    Port_to_coord := Port_to(sender_of(Reconfig));
     $\forall p \in (\text{Set\_of\_Ports} - \text{Port\_to\_coord})$ 
        send Reconfig(node_list + my_id, frag_id)
        through p.

```

### 3.4 Detection of potential cycles

Since a reconfiguration wave propagates through the network by message flooding, the Reconfig(*node\_list*, *frag\_id*) message received by *S'* may be due to cycles in the physical topology, which is detected and handled as below:

1. The Reconfig message is a copy of the message received by *S'* earlier, which may arise due to the sending of this message by some node *S''* through its ports. A potential cycle that may exist in this case is indicated by the condition *Coord\_so\_far* = *frag\_id*. If the Reconfig message is received over a graph edge, *S'* infers that the message has entered the fragment to which *S'* originally belonged at two different points resulting in the *frag\_id* colliding with itself over the edge. So *S'* needs to resolve the contention so that its fragment gets connected to *frag\_id* at one point only. If the Reconfig message is not received over a graph edge, *S'* responds with a No\_contention.
2. The Reconfig message is a copy of the message sent by *S'* earlier through a port but now received through another port. This message loop is indicated by the condition *my\_id*(*S'*)  $\in$  *node\_list*. In this case, *S'* sends a No\_contention through sender\_of(Reconfig).

The actions are given by the following code skeleton:

```

receive Reconfig(..., frag_id)
if (status = WAIT)
    e := Port_to(sender_of(Reconfig));
    if (frag_id = Coord_so_far && e  $\notin$  Get_Port)
        send No_contention through e;
        exit;
    if (my_id  $\in$  node_list)
        send No_contention through e;
        exit;
    resolve contention between Coord_so_far
    and frag_id.

```

Consider the nodes  $\{S_3, S_4, S_5, S_6, S_7, S_9\}$  in the topology illustrated in Figure 2. Suppose only *S*<sub>6</sub> detects the failure of *S*<sub>9</sub>. *S*<sub>6</sub> sends a Reconfig(*id*(*S*<sub>6</sub>), {*id*(*S*<sub>6</sub>)}) message to *S*<sub>4</sub> and *S*<sub>5</sub>. the message arriving at *S*<sub>4</sub> is forwarded to *S*<sub>3</sub> and *S*<sub>7</sub> over edges, while that at *S*<sub>5</sub> is forwarded to *S*<sub>3</sub> over a link. When *S*<sub>7</sub> receives the Reconfig as forwarded from *S*<sub>3</sub> over a link after that from *S*<sub>4</sub>, *S*<sub>7</sub> responds to *S*<sub>3</sub>

with a No\_contention. The latter prevents creation of an edge between *S*<sub>3</sub> and *S*<sub>7</sub>, thereby avoiding a potential cycle between  $\{S_3, S_4, S_7\}$ . If *S*<sub>3</sub> received the Reconfig from *S*<sub>4</sub> and forwarded it to *S*<sub>5</sub>, then *S*<sub>3</sub> and *S*<sub>5</sub> send a No\_contention to each other. If *S*<sub>3</sub> received the Reconfig from *S*<sub>5</sub> and forwarded it to *S*<sub>4</sub>, a contention occurs between *S*<sub>3</sub> and *S*<sub>4</sub> over an edge which, upon resolution, allows either *S*<sub>3</sub> to win generating a No\_contention from *S*<sub>4</sub> to *S*<sub>6</sub> or *S*<sub>4</sub> to win generating a No\_contention from *S*<sub>3</sub> to *S*<sub>5</sub>. The case of *S*<sub>6</sub> receiving a Reconfig over the path [*S*<sub>5</sub>, *S*<sub>3</sub>, *S*<sub>4</sub>] is a message loop, detectable from the *node\_list* = {*id*(*S*<sub>6</sub>), *id*(*S*<sub>5</sub>), *id*(*S*<sub>3</sub>), *id*(*S*<sub>4</sub>)}.

Upon a contention resolution, *Port\_to\_coord* is flipped in the direction towards the winning fragment and *Coord\_so\_far* is updated to indicate this fragment id.

### 3.5 Fragment growth

Suppose a node *S* has sent out a Reconfig(..., *F*<sub>*i*</sub>). Upon receipt of No\_contention messages through all the ports where a response is awaited for, *S* can infer that there is no contention in the downstream direction. In this case, if there is a fragment rooted at *S*, then *S* should connect its vertex to a fragment in the upstream direction. For this purpose, *S* sends an Accepted message through *Port\_to\_coord*.

Thus a node *S* may receive either an Accepted message or a No\_contention message through a port *p*, with the Accepted message indicating a successful acquisition of the fragments accessible through *p*. The termination condition for the algorithm is thus:

$$\text{Set\_of\_Ports} = \text{Port\_to\_coord} + Q,$$

where *Q* is the set of ports through which either an Accepted or a No\_contention message has been received. Upon evaluation of the termination condition to true, *S* sends:

- An Accepted message through the *Port\_to\_coord* if: i) *S* has received an Accepted message through at least one port, or ii) a No\_contention message has arrived through a port containing an edge and *Port\_to\_coord* does not contain an edge;
- A No\_contention message through the *Port\_to\_coord* if *S* has received No\_contention messages through all the ports.

In either case, *S* returns to the IDLE state completing its part of the reconfiguration. Both the messages inform the neighbor *S'* linked through *Port\_to\_coord* that *S* has accepted *F*<sub>*i*</sub> as the coordinator. Suppose *S* sends an Accepted message. If there is no edge over the link between *S* and *S'*, an edge is created by these nodes by invoking Assign\_edge(*Port\_to\_coord*) and Assign\_edge(Port\_to(sender\_of(Accepted))) respectively. Such an edge interconnects the fragment with a vertex at *S* and that at *S'*. Thus the Accepted message allows the fragments acquired by *S* to be handed over to *S'*. These actions are illustrated by the following code skeleton:

```

receive Accepted/No_contention
  if  $\forall p \in (\text{Set\_of\_ports} - \text{Port\_to\_coord})$ 
    if (at least one response received is Accepted)
      send Accepted through Port_to_coord;
    if Port_to_coord  $\notin$  Get_Port
      Assign_Edge(Port_to_coord);
    else /* All responses received are No_contention */
      if  $\exists q \in (\text{Set\_of\_ports} - \text{Port\_to\_coord})$ 
        such that  $q \in \text{Get\_Port} \ \&\&$ 
          Port_to_coord  $\notin$  Get_Port
            send Accepted through Port_to_coord;
            Assign_Edge(Port_to_coord);
      else
        send No_contention through Port_to_coord.

```

For this purpose, the *recd\_reply*[ ] also indicates whether an Accepted or a No\_contention was received through a given port.

Before an Accepted message reaches  $S'$  it is possible that  $S'$  might have changed the *Coord\_so\_far* and consequently the *Port\_to\_coord* to a wave originated by a fragment  $F_j$ . Thus the fragment grown by the wave from  $F_i$  is inherited by the wave from  $F_j$ . This can be viewed as fragment acquisition by  $F_j$  from  $F_i$ .

As can be seen, various fragments keep combining with each other until the algorithm terminates at the elected coordinator node where *Port\_to\_coord* = NULL. Since the edge created by an Accepted message is never deleted, a newly created fragment is always incremental to the previous fragments. These aspects guarantee the steady growth of fragments. The propagation of information across various fragments is an instance of diffusing computations [7].

### 3.6 A sample scenario

Consider the fragmentation scenario illustrated in Figure 2. Suppose  $S_{10}$  is the only initiator of reconfiguration.  $S_{10}$  will send a Reconfig message to  $S_6$  and  $S_{12}$ . The message through  $S_6$  will eventually reach  $S_{13}$  through  $S_8$ , whereupon  $S_{13}$  generates a No\_contention message to  $S_8$ . Since  $S_8$  is a vertex in the graph and has only a link to  $S_7$ ,  $S_8$  transforms the No\_contention message into an Accepted message and sends it to  $S_7$ , which creates an edge connecting  $[S_7, S_8]$ . Suppose the Accepted message passes through  $[S_4, S_3, S_1]$ . When the message traverses from  $S_1$  to  $S_2$ , an edge is created connecting  $S_1$  and  $S_2$ . The Accepted message then passes through  $[S_2, S_5, S_6]$  and upon traversing from  $S_6$  to  $S_{10}$  causes the latter nodes to be connected by an edge. The Reconfig message which reaches  $S_{11}$  causes an Accepted message to be generated since  $S_{11}$  is a vertex and has only a link to  $S_{12}$ . This Accepted message, upon receipt by  $S_{12}$ , creates an edge connecting  $S_{11}$  and  $S_{12}$  and then gets forwarded to  $S_{10}$ . When both the Accepted messages arrive at  $S_{10}$ , the algorithm terminates since  $S_{10}$  has its *Port\_to\_coord* set to NULL.

Thus, we may have an algorithm which, using a contention resolution mechanism, is able to reconfigure a

graph. In the following sections 4 and 5, we discuss contention resolution using respectively: (i) ranking among fragments, and (ii) random arbitration among fragments.

## 4 Ranking based contention resolution

In this section, we will use unique ranking among node ids as a mechanism to resolve contention between two fragments. The network guarantees that the ranking remains stable during an execution of the algorithm.

### 4.1 Winning contention

We say a fragment  $F_i$  is ranked higher than a fragment  $F_j$  if  $\text{id}(F_i) > \text{id}(F_j)$ , and vice-versa. Consider the meeting of fragments  $F_i$  and  $F_j$  when they attempt to expand. If  $F_i$  is ranked higher than  $F_j$ , then  $F_i$  is the winner and  $F_j$  is the loser; otherwise  $F_j$  is the winner and  $F_i$  is the loser.

Suppose a node  $S$  is in the WAIT state, has *Coord\_so\_far* =  $F_j$ , and receives a Reconfig( $\dots, F_i$ ). The following code skeleton illustrates the activities of  $S$ :

```

if ( (Coord_so_far >  $F_i$ ) ||
    (Coord_so_far =  $F_i$  &&
     my_id >  $\text{id}(\text{sender\_of}(\text{Reconfig}))$ ) )
  send Stop(Coord_so_far) through
    Port_to( $\text{sender\_of}(\text{Reconfig})$ );
else
  Coord_so_far :=  $F_i$ ;
  send Stop( $F_i$ ) through Port_to_coord;
  Port_to_coord := Port_to( $\text{sender\_of}(\text{Reconfig})$ );

```

As can be seen, when  $F_j$  wins,  $S$  sends a Stop message to the sender of Reconfig to inform all vertices in  $F_i$  that no other fragment ranked less than  $F_j$  will be contending on the path beyond the point at which the collision occurred. When  $F_j$  loses,  $S$  sends a Stop through *Port\_to\_coord* to indicate to other vertices in  $F_j$  about its losing the contention and then flips its *Port\_to\_coord* towards the winning fragment  $F_i$ .  $S$  then sets *coord\_so\_far* to  $F_i$ .

### 4.2 Fragment acquisition

Consider the arrival of Stop( $F_i$ ) at a node  $S$  with *coord\_so\_far* =  $F_j$ . Suppose  $F_i$  is of higher rank than  $F_j$ . Then  $S$  sends Stop( $F_i$ ) through *Port\_to\_coord* to notify the win of  $F_i$  over  $F_j$  to all vertices in  $F_j$ .  $S$  then flips the *Port\_to\_coord* towards the sender of Stop( $F_i$ ). If  $F_i$  is of lower rank than  $F_j$  the  $S$  sends a Stop( $F_j$ ) to the sender of Stop( $F_i$ ) to notify the win of  $F_j$  over  $F_i$  to all vertices in  $F_i$ . This is given by the following code skeleton:

```

receive Stop( $F_i$ ) through a port  $p$ 
  if ( $F_i$  > coord_so_far)
    coord_so_far :=  $F_i$ ;
    send Stop( $F_i$ ) through Port_to_coord;
    Port_to_coord :=  $p$ ;
  if ( $F_i$  = coord_so_far)
    if (Port_to_coord  $\notin$  Get_Port)
      send No_contention through Port_to_coord;

```

```

    mark recd_reply[Port_to_coord] as having
    received No_contention;
  else
    send Stop( $F_i$ ) through Port_to_coord;
    Port_to_coord :=  $p$ ;
    if ( $F_i < \text{coord\_so\_far}$ )
      send Stop( $F_j$ ) through  $p$ .

```

Since *coord\_so\_far* is updated to  $F_i$  only when its current value (i.e.,  $\text{id}(F_j)$ ) is less than  $\text{id}(F_i)$  and it is accompanied by a **Stop**( $F_i$ ) to all vertices in  $F_j$ , each of these vertices will eventually agree at the winning of  $F_i$  by setting their *Coord\_so\_far* to  $\text{id}(F_i)$ . Thus the fragment  $F_j$  is acquired by  $F_i$  and the winning of  $F_i$  is propagated across all vertices in  $F_j$ . As can be seen, this guarantees monotonically increasing values of *Coord\_so\_far*. It means that a tree is rooted currently at some vertex acting as the coordinator, and can change to any other vertex with a higher rank.

#### 4.3 Deterministic Termination

Use of ranking to arbitrate between colliding fragments terminates the algorithm deterministically. Suppose the fragments  $F_1$ ,  $F_2$ , and  $F_4$  start the reconfiguration and  $F_3$  is too slow to start the reconfiguration. By the time  $F_3$  starts, some other fragment subdues it. Now it is up to  $F_1$ ,  $F_2$ , and  $F_4$  to complete the reconfiguration. Some of the edges could be accumulated by  $F_1$ , some by  $F_2$  and some by  $F_4$ . Let us say  $\text{id}(F_4) > \text{id}(F_2) > \text{id}(F_1)$ . Since this information is stable, the protocol will eventually terminate with  $F_4$  as the coordinator. This is because whenever there is contention, the higher ranked wave wins. Thus  $F_2$  can win over  $F_1$ , but  $F_4$  will win over both  $F_2$  and  $F_1$ . This happens in finite time since message transmission delay over a link is bounded and contentions will be resolved within the maximum diameter of the network  $\mathcal{D}$ .

### 5 Random Arbitration to resolve contention

When the ranking information is not available to the algorithm, it is not possible to declare a fragment  $F_i$  as a winner over a fragment  $F_j$  or vice-versa, when the two fragments meet resulting in a collision. Random arbitration can declare any of the fragments as a winner in a probabilistic manner. This however leads to non-determinism, because  $F_i$  could win over a fragment  $F_j$  at one meeting point but could loose at another point. So this non-determinism has to be tackled in the algorithm.

#### 5.1 Random backoff

When two fragments  $F_i$  and  $F_j$  meet, resulting in a collision, the algorithm requires both the fragments to backoff for a random period of time before attempting reconfiguration in the direction where the collision occurred. This may allow one fragment to win over the other after one or more attempts.

In the algorithm, the two nodes detecting the collision move from the WAIT state into a BACKOFF state. In this state, each node will retain its *Coord\_so\_far* and *Port\_to\_coord* values, but remains inactive (or frozen), i.e., does not send any messages for a certain interval of time. After this interval, a node say  $F_i$  re-sends the **Reconfig** message and moves into WAIT state. If the message reaches  $F_j$  while the latter is still in the BACKOFF state, then  $F_i$  is declared a winner over  $F_j$ . But if  $F_j$  had become active in the meanwhile, there will be contention again, and the same backoff procedure will be repeated. A winner then finally continues the reconfiguration.

Consider the arrival of **Reconfig**( $\dots, F_i$ ) at a node  $S$  in a WAIT state and *Coord\_so\_far* =  $F_j$ . Contention is said to have occurred if  $F_i \neq F_j$ . This causes the wave  $F_j$  to backoff. Since the message **Reconfig**( $\dots, F_j$ ) sent by  $S$  will also suffer collision at the neighbor that sent the **Reconfig**( $\dots, F_i$ ),  $S$  simply assumes that it has got the latter fragment back and that it is going to restart the reconfiguration after the backoff interval. Also, when  $S$  moves to the BACKOFF state, it sends an **Abort** message through all ports over which neither an **Accepted** nor a **No\_contention** has been received. The **Abort** message informs all vertices in the fragment about the contention. This is illustrated by the following code skeleton:

```

receive Reconfig( $\dots, F_i$ )
  if (status = WAIT)
    if ( (Coord_so_far  $\neq F_i$ ) || (Coord_so_far =  $F_i$  &&
      Port_To(sender_of(Reconfig))  $\in$  Get_Port))
      status := BACKOFF;
       $\forall p \in (\text{Set\_of\_ports} - \text{sender\_of}(\text{Reconfig}))$ 
        if recd_reply[ $p$ ] indicates no response has
          been received
            send Abort through  $p$ ;
      start random backoff timer.

```

The **Abort** message arriving at a node  $S$  in a WAIT state freezes its reconfiguration activity by placing  $S$  in BACKOFF state. The reconfiguration at  $S$  is re-activated from where it was frozen upon receiving a **Reconfig** message again. However, if  $S$  is in IDLE or BACKOFF state, the arrival of an **Abort** does not affect  $S$ . This is illustrated by the following code skeleton:

```

receive Abort
  if (status  $\neq$  BACKOFF && status  $\neq$  IDLE)
    status := BACKOFF;
    if (coord_so_far  $\neq$  my_id)
       $\forall p \in (\text{Set\_of\_ports} - \text{sender\_of}(\text{Abort}))$ 
        if recd_reply[ $p$ ] indicates no response has
          been received
            send Abort through  $p$ .

```

As mentioned earlier, a node  $S$  in the BACKOFF state does not send any messages. However, if an **Accepted** or a **No\_contention** is received on a port while in the BACKOFF state, the *recd\_reply*[ $p$ ] is updated accordingly.




Figure 3: Message Complexity Vs Time Complexity

## 5.2 Restart of reconfiguration

The BACKOFF state can be seen as a frozen state because the node, having generated **Abort** messages, wakes up upon arrival of a **Reconfig** message. In addition, a node in BACKOFF state may also wake up upon expiry of the backoff timer. If a node in BACKOFF state is woken up by the arrival of a **Reconfig** message, the node loses the contention in favor of the fragment that sent the **Reconfig**, and hence cancels the backoff timer. Upon waking up, a node re-sends a reconfiguration message on all ports over which neither an **Accepted** nor a **No\_contention** was received. This is illustrated by the following code skeleton:

```

receive Reconfig(...,  $F_i$ )
  if (status = BACKOFF)
    cancel backoff timer;
    status := WAIT;
    if (Coord_so_far =  $F_i$ )
      if  $\forall p \in (\text{Set\_of\_ports} - \text{sender\_of}(\text{Reconfig}))$ 
        recd_reply[p] indicates no response has
        been received
        send Reconfig(..., Coord_so_far) through p.

```

The **Accepted** or **No\_contention** messages received by a node while in backoff state cause the edges created along the direction of these messages to be frozen for acquisition by a subsequently winning wave.

Consider, for instance, the fragments  $F_i$  and  $F_j$  contending at two different points.  $F_i$  may win over  $F_j$  at one point and lose at the other point. In this case, both  $F_i$  and  $F_j$  will regenerate the reconfiguration activities at the winning points.  $F_i$  and  $F_j$  may collide again in these regenerated activities, possibly repeating the above scenario. This continues until one of the fragments wins at all the points. Since the number of times a fragment will backoff is unknown, we cannot place an upper bound on the execution time of the protocol. However, since every reconfiguration activity potentially causes addition of new edges and the already acquired edges are never deleted, the graph grows steadily and finally terminates at the fragment which emerged as the eventual winner.




Figure 4: Effect of varying Backoff time on the Message Complexity

## 6 Evaluation of the algorithms

We discuss in this section the performance of ranking based and random backoff based reconfiguration algorithm by studying their simulated behavior.

### 6.1 Simulation Studies

The reconfiguration algorithms are studied by employing discrete event simulation. The message transfer delay over a link is modelled as a random variable. For fiber optic networks, the link delay can range from 5-10 microseconds for a 100 byte control message. Various arbitrary network topologies were generated by specifying the nodes and the links between them and the initial graph.

### 6.2 Performance evaluation

As the density of the graph increased, i.e., as the average degree of connectivity increased, the time and message complexity increased. For sparse graphs, fewer messages were generated as compared to graphs which were more dense, as shown in figure 3. Here, for the ranking approach, as we increased the average degree of connectivity  $\bar{e}$ , the number of messages increased within the range 20 - 70. Correspondingly, the time taken to termination of the protocol also increased within a range of 12 - 22 time units. For the Random backoff approach, the messages required increased on increasing the average degree of connectivity  $\bar{e}$  as well as by decreasing the random backoff interval.

The time unit used in the reconfiguration latency represents the average transfer delay over a link. Thus a time delay of 100 units to reconfigure means a reconfiguration latency of 1 millisecond (approx) on a high speed network. We believe that such a latency will have a negligible effect on the network applications. In a video conferencing for instance, a glitch in the workstation window may not last for more than a few milliseconds.

Using the random arbitration technique, we observed that there was a significant change in the number of messages required, as shown in figure 4. Here, for a graph G1, with a lower average degree of connectivity  $\bar{e}$ , the number of messages decreased as the backoff time was increased to 3 time units. Beyond that, the number of messages again




Figure 5: Effect of varying Backoff time on the Time Complexity




Figure 6: Probability distribution of reconfiguration delay for random backoff approach

started increasing. This shows that a certain backoff value is optimum for a certain graph. For graphs G1 and G2, where  $\bar{e}(G1) < \bar{e}(G2) < \bar{e}(G3)$ , as  $\bar{e}$  was increased, the number of messages increased.

Similarly, there was a significant change in the time required to termination, as shown in figure 5. Here, as the backoff time was varied between 2 and 4 time units, the total time taken varied between 10 and 20 time units. As the reconfiguration protocol terminates, we observed that the reconfigured path was incremental.

The probabilistic termination of the random backoff approach can be seen in figure 6. Here we simulated the algorithm for various random number sequences and upon varying the backoff time, the total time varied within an acceptable range.

## 7 Summary

The paper described reconfiguration of graph-connected paths in a network in the presence of node failures. We modeled the reconfiguration problem as one of connecting two or more subgraphs together which hitherto were not connected with one another. The fragment interconnection require coordination among the concerned fragments so that the new graph composed from the fragments does not contain cycles. The coordination problem manifests itself as resolving contention among the fragments to determine the edges to be created. To deal with environments

with potentially large number of fragments, the paper described two classes of protocols:

- Based on a pre-established ranking among the fragments;
- Based on a random arbitration among the fragments.

The paper designed the reconfiguration algorithms and evaluated them. We are currently extending the algorithms to a general case of dealing with multiple failures and joining/leaving of nodes [6]. For instance, dealing with multiple failures requires composing execution of multiple instances of the reconfiguration algorithms discussed above.

The reconfiguration algorithms proposed in this paper are different from that in Autonet [5] in that this work requires networkwide update of the entire physical topology information whenever a change occurs. [8] also proposed a reconfiguration algorithm but it requires some centralized information ('replacement sets') to reconfigure the tree. [9] gives an algorithm for constructing a spanning tree but does not deal with reconfigurations.

## References

- [1] B. M. Waxman. **Routing of Multipoint Connections**. In *IEEE Journal on Selected Areas in Communications*, Vol.SAC-6, No.9, pp.1617-1622, Dec 1988.
- [2] K. Ravindran, M. Sankhla and P. Gupta. **Multicast Models and Routing Algorithms for High Speed Multi-service Networks** *Technical Report*, Kansas State University, Feb 1992.
- [3] R. Perlman. **An Algorithm for Distributed Computation of a Spanning Tree in an Extended LAN**. In *Proc. Communication Architectures and Protocols*, ACM SIGCOMM, pp. 44-53, Sept 1985.
- [4] D. Bertsekas and R. Gallager. **Routing in Data Networks** Chapter 5 in *Routing in Data Networks*, Prentice Hall Publ. Co., 1992.
- [5] M. D. Schroeder and T. L. Rodeheffer. **Automatic Reconfiguration in Autonet** In *Technical Report*, Digital Equipment Corporation (Systems Research Center), Palo Alto (CA), 1991.
- [6] K. Ravindran and G. Singh. **Algorithmic Frameworks for Reconfiguration of Spanning Trees in Distributed networks**. *Work in Progress*, Kansas State University, Oct. 1992.
- [7] E. Dijkstra and C. Scholte. **Termination Detection for Diffusing Computations** *Information Processing Letters*, 11(1), Aug 1980.
- [8] I. A. Cimet, C. Cheng and S. P.R. Kumar. **On the Design of Resilient Protocols for Spanning Tree Problems**. In *IEEE transactions in Communications* 1989.
- [9] R. G. Gallager, P. Humblet and P. Spira. **A distributed algorithm for minimal spanning tree**, *ACM Transactions of Programming languages and Systems*, 30(12), Dec 1983.